

CS 4240: Compilers and Interpreters

Phase 3: Register allocation and Instruction selection and generation

Total Points: 125

Due Date: April 20, 2014

This phase of the project consists of two parts plus a bonus part. You will be allocating registers for the IR stream you produced in phase 2, and you will be selecting the final MIPS instructions for the instruction stream. For bonus points, you can implement function calls. You will be able to run your output on a MIPS simulator by the end of this phase.

Part 1: Register Allocation

Your stream of IR instructions from phase 2 relied on an unlimited number of temporaries to store intermediate results. In this part, the goal is to remove these temporaries and achieve an instruction stream that is faithful to MIPS' register set. Wikipedia's page on MIPS includes details on all of the registers you will need for this part:

http://en.wikipedia.org/wiki/MIPS_architecture

The general idea of register allocation is straightforward. At any given point, the processor can only hold a finite number of values in its register file, and the values being manipulated need to be in those registers. If there are not enough registers, we have *register spill*, and extra store and load instructions are needed. The way to handle register spill is to use memory. Because a *live variable* will be used in the near future, it must not only be stored in memory but also loaded back into the register file before its use. All variables are allocated space during compile-time in the .data memory segment (see part 2). When a variable must be brought into a register, its memory address in this segment is known. Likewise, when it must be stored, it will be stored at this known address.

There is a dense body of knowledge surrounding register allocation, and it is a critical compiler phase for good performance for obvious reasons. For the purposes of this project, we are interested in correctness but not high levels of performance. You will implement three different register allocation schemes.

Naive

The most naive allocation scheme is one in which there is no analysis. Before each instruction, its operands are loaded into registers; the instruction then executes; and finally the result is stored back into that variable's home location in memory. Thus, for each instruction in the IR

stream, you will generate and insert the necessary load(s) before that instruction, and you will generate and insert the necessary store(s) after that instruction. This scheme is the slowest, but it will produce correct, working code.


CFG Construction and Intra-block Allocation

An improvement on the naive scheme is to locate the *basic blocks* in the stream, construct a *control flow graph* (CFG) of those blocks, and carry out *graph coloring* at the intra-block level (i.e. only within each basic block). We have provided the resources for how to detect basic blocks, build the control flow graph, and perform graph coloring on each basic block. Notice that at the start of each block, you will need to load a set of variables that you expect to use, and similarly, at the end of each block, you will need to store all values from your registers to memory.

EBB Construction and Intra-EBB Allocation

Extended basic blocks (EBBs) allow us to perform register allocation on a larger section of code (thereby allowing for better register allocation). An extended basic block is a collection of connected basic blocks such that only the first basic block may have more than one predecessor in the control flow graph. For illustration, consider Figure 1. Each dashed box encloses an EBB. A complete algorithm for constructing all EBBs is given in Algorithm 1. (Source: *Advanced Compiler Design Implementation* by Muchnick)

Figure 1: Control flow graph with extended basic blocks enclosed in dashed boxes



Algorithm 1: Construction of all EBBs

```
entry: Node
EbbRoots: set of Node
AllEbbs: set of (Node × set of Node)

procedure main()
begin
    Build_All_Ebbs(entry, Succ, Pred)
end

procedure Build_All_Ebbs(r, Succ, Pred)
    r: in Node
    Succ, Pred: in Node → set of Node
begin
    x: Node
    s: Node × set of Node
    EbbRoots := {r}
    AllEbbs := ∅

    while EbbRoots ≠ ∅ do
        x := ♦ EbbRoots
        EbbRoots -= {x}
        if ∀s ∈ AllEbbs (s@1 ≠ x) then
            AllEbbs U= {<x, Build_Ebb(x, Succ, Pred)>}
        fi
    od
end

procedure Build_Ebb(r, Succ, Pred) returns set of Node
    r: in Node
    Succ, Pred: in Node → set of Node
begin
    Ebb := ∅: set of Node
    Add_Bbs(r, Ebb, Succ, Pred)
    return Ebb
end
```

```

procedure Add_Bbs(r, Ebb, Succ, Pred)
    r: in Node
    Ebb: inout set of Node
    Succ, Pred: in Node → set of Node
begin
    x: Node
    Ebb U= {r}
    for each x ∈ Succ(r) do
        if |Pred(x)| = 1 & x ∉ Ebb then
            Add_Bbs(x, Ebb, Succ, Pred)
        elif x ∉ EbbRoots then
            EbbRoots U= {x}
        fi
    od
end

```

Once all EBBs have been constructed, you will allocate registers within an EBB using the same graph coloring technique from the previous scheme. Similarly, you will load needed values at the start of an EBB and store values at the end of the EBB.

Your register allocation implementations should take correct IR code as input and produce modified IR code for use in part 2. We do not need to see the output of this part, but it may be useful to produce an output file for debug purposes.

Part 2: Instruction Selection and Code Generation

You only have to write a single implementation of instruction selection and code generation. Any of the three implementations from part 1 should produce IR code with correct register allocation, and this stream will be used as input to your part 2 solution.

The instruction selection is a matter of converting the IR code you've generated to the appropriate MIPS assembly code. The IR code from phase 2 is actually a relatively close match to the MIPS code you are expected to generate (with some exceptions). Wikipedia's page on the MIPS architecture includes the assembly instruction supported by MIPS. Another challenge is ensuring the .data segment is created correctly. This link may help to understand this:

<http://www.cs.umd.edu/class/sum2003/cmsc311/Notes/Mips/dataseq.html>

You will be running your generated MIPS code on a simulator called SPIM (see the end of this document). For your code to run properly on SPIM, you should generate a "main" label where your program's statement sequences begin (i.e. just after the "in" of your program's "let-in-end").

You must also generate a final instruction, “jr \$ra”, to return to the caller of your program. Supplied examples will help demonstrate this.

For this part, you may safely ignore direct support in the language for string comparisons. Also, unless you are implementing function calls for bonus points, you may safely ignore instruction selection for function calls, as well.

This part of the project should take the modified code produced in part 1 and output correct assembly code. Your final assembly should run on SPIM.

Part 3 (bonus): Function Calls

If you choose to support function calls for bonus points, then hack on! There are of course many resources on the internet on MIPS calling conventions. This link is a good start:

<http://people.cs.pitt.edu/~xujie/cs447/Mips/sub.html>

Turn-in

Grading

1. Register allocation code (60 points)
 - a. Naive (20 points)
 - b. CFG and intra-block allocation (25 points)
 - c. EBB and intra-EBB allocation (15 points)
2. Instruction selection and generation code (35 points)
3. Passes tests using generated code executing on simulator (20 points)
4. Report (design internals, how to build, run, etc.) (10 points)

Bonus:

1. Function calls (working MIPS code) (40 points)

SPIM

SPIM is a MIPS simulator you will use to run your compiled, MIPS assembly programs. To download SPIM, use the following link:

<http://spimsimulator.sourceforge.net/>

Once SPIM is running, execute the hello-world example:

File > Load File

Select <path-to-spim-simulator>/helloworld.s
Click the green arrow to run it
SPIM console should display “Hello World”

Troubleshooting:

If you run into trouble with their binary package, you can download the source with svn:
svn checkout svn://svn.code.sf.net/p/spimsimulator/code/ spimsimulator-code

Then you can build the source with QtCreator. Use this to get the most recent Qt:
<https://qt-project.org/downloads>

After running Qt-Creator, open the project:
<path-to-spim-simulator>/QtSpim/QtSpim.pro.

Build and run the project. This should launch SPIM, and then you can proceed with the hello-world example.